

92586 Computational Linguistics

Lesson 14. Hands on Word Embeddings

Alberto Barrón-Cedeño

Alma Mater Studiorum-Università di Bologna
a.barron@unibo.it @albarron_

22/04/2021



Previously

- Skip-gram
- CBOW

Table of Contents

Pre-Trained Models

Gensim

Model Construction

GloVe

fastText

Chapter 6 of Lane et al. (2019)

Pre-Trained Models

Some Pre-Trained Models

Model	Provider	Description
word2vec	Google	300D from English Google News articles ¹
fastText	Facebook	157 languages from Wikipedia and Crawl ²
word2vec/GloVe	CNR	Italian embeddings from the Wikipedia
word2vec	UCampania	Italian embeddings ³

There are many pre-trained models and diverse libraries to handle them.

Just go to your favorite search engine

¹<https://drive.google.com/file/d/0B7XkCwpI5KDYN1NUTT1SS21pQmM>

²<https://fasttext.cc>


³<https://mlunicampania.gitlab.io/italian-word2vec/>

Gensim

Gensim

Gensim

- ▶ Scalable, open source, and efficient Python library
- ▶ It includes many resources, including word2vec, doc2vec, FastText, LDA, and more
- ▶ All information, including very nice manuals at <https://radimrehurek.com/gensim/>

 Let us see

Gensim

Most similar items


```
word_vectors.most_similar()
```

Among the most interesting parameters:

positive list of vectors to be added together before looking for the neighbours

negative subtraction (or exclusion) of the elements

topn number of elements to retrieve


 Let us see

Gensim

Least similar items (closed set)

```
word_vectors.doesnt_match()
```


It returns the element from the input list with the lowest similarity wrt the rest

 Let us see

Gensim


More operations

Adding and Subtracting We can use `most_similar()` again, this time with the negative parameter

 Let us see

Computing similarities


```
word_vectors.similarity()
```

 Let us see

Gensim

Getting the Vectors

Gensim (and other libraries) have coded these interfaces to perform operations, but one might want to go beyond `word_vectors[word]`

 Let us see

Model Construction

Model Construction

Considerations

- ▶ If you are working in other language than English, Google's provided word2vec is not an option (FastText might be)
- ▶ Google's word2vec is built on news; fastText is built on the Wikipedia. . . analysing scientific papers or literature?

Probably not

- ▶ You want to work on COVID-19 or any other recent topic?
Many relevant terms **wont appear**

Alternatives

- ▶ Opting for some of the previous representations
- ▶ **Build your own model**

Model Construction

Pre-Processing

Typical pre-processing pipeline

- ▶ Tokenization
- ▶ Lowercasing (optional)
- ▶ Sentence splitting

Input Embedded list of tokenised sentences

$$[[w_{0,0} \ w_{0,1} \ w_{0,2} \ \dots \ w_{0,k}], [w_{1,0} \ w_{1,1} \ w_{1,2} \ \dots \ w_{1,l}], \dots [w_{x,0} \ w_{x,1} \ \dots \ w_{x,m}]]$$

Model Construction

Training

Training the word2vec model with gensim

Documentation:

<https://radimrehurek.com/gensim/models/word2vec>

 Let us see

Considerations

- ▶ A few minutes are necessary for small corpora
(Brown took me 2 minutes on a 2.5GHz Quad-Core i7, 16GB RAM)
- ▶ Large corpora (e.g., the Wikipedia) can require a significant amount of time/memory

Model Construction


Trimming and Saving

Reminder We do not care about the output

`model.init_sims(replace=True)`

- ▶ Freezes the model
- ▶ Stores the hidden-layer weights
- ▶ Discards the output-layer weights

Now we simply have to save the model with `model.save()`

 Let us see

GloVe

GloVe

Global Vectors Pennington et al. (2014)⁴

- ▶ It uses a global word-word co-occurrence matrix
- ▶ Learning objective: word vectors such that their dot product equals the logarithm of the words' probability of co-occurrence
- ▶ It produces similar matrices to word2vec
- ▶ It converges, even with smaller corpora
- ▶ It is more accurate with the same amount of data

⁴<https://nlp.stanford.edu/projects/glove/>

GloVe

GloVe vs word2vec

RaRe Technologies comparison⁵

Settings: 600 dims • context window of 10 • 1.9B words of en Wikipedia.

Algorithm	acc (word analogy)	wallclock time	peak RAM (MB)
I/O only	–	3m	25
GloVe, 10 epochs, lr 0.05	67.1	4h12m	9,414
GloVe, 100 epochs, lr 0.05	67.3	18h39m	9,452
word2vec, hierarchical skipgram, 1 epoch	57.4	3h10m	266
word2vec, negative sampling (10 samples), 1 epoch	68.3	8h38m	628
word2vec, Google 300d	55.3	–	–

⁵rare-technologies.com/making-sense-of-Word2vec/#glove_vs_word2vec

fastText

fastText

Predicts the surrounding **character** [2, 3]-**grams** rather than the surrounding words Bojanowski et al. (2017)⁶

- Pre-trained models available in 250+ languages
- Built on Wikipedia editions (variable quality)

Models available at <https://github.com/facebookresearch/fastText/blob/master/docs/crawl-vectors.md>

Example:

```
wget -c \
https://dl.fbaipublicfiles.com/fasttext/vectors-crawl/cc.it.300.bin.gz
```

 Let us see

⁶<https://github.com/facebookresearch/fastText>

Some Remarks

LSA is a better (faster) option for long documents e.g., for clustering

Online learning An existing model can be *adapted* (but no new words can be added)

doc2vec possible representation based on linear combinations of word2vec

References

Bojanowski, P., E. Grave, A. Joulin, and T. Mikolov
2017. Enriching Word Vectors with Subword Information.
Transactions of the Association for Computational Linguistics,
5:135–146.

Lane, H., C. Howard, and H. Hapkem
2019. *Natural Language Processing in Action*. Shelter Island,
NY: Manning Publication Co.

Pennington, J., R. Socher, and C. Manning
2014. GloVe: Global Vectors for Word Representation. In
Empirical Methods in Natural Language Processing (EMNLP),
Pp. 1532–1543.