

# 92586 Computational Linguistics

## Lesson 16. From document representations, towards sequences

Alberto Barrón-Cedeño

Alma Mater Studiorum-Università di Bologna  
a.barron@unibo.it      @\_albarron\_

23/04/2021



# Previously

- Training and loading (existing) embeddings
- Visualisation

# Table of Contents

1 Doc2vec

2 Prologue to CNN and RNN

3 CNN

Chapters 6 and 7 of Lane et al. (2019)

# **Doc2vec**

# Doc2vec

**Objective** Computing a vectorial representation of a document.

# Doc2vec

**Objective** Computing a vectorial representation of a document.

Same idea as with word2vec: a NN to predict words

# Doc2vec

**Objective** Computing a vectorial representation of a document.

Same idea as with word2vec: a NN to predict words

## Input

- $k$  context words (optional)
- A unique ID of the sentence/paragraph/document

# Doc2vec

**Objective** Computing a vectorial representation of a document.

Same idea as with word2vec: a NN to predict words

## Input

- $k$  context words (optional)
- A unique ID of the sentence/paragraph/document

## Output

- 1 target word

# Doc2vec

**Objective** Computing a vectorial representation of a document.

Same idea as with word2vec: a NN to predict words

## Input

- $k$  context words (optional)
- A unique ID of the sentence/paragraph/document

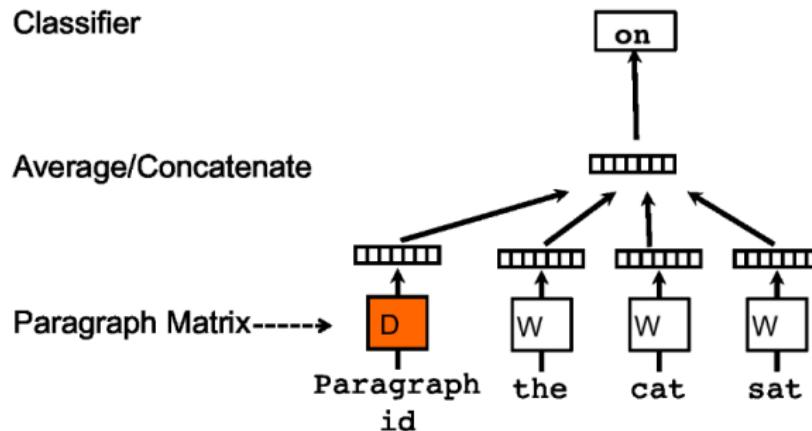
## Output

- 1 target word
- The paragraph vector is unique among all documents
- The word vectors are shared among all documents
- The document vector is computed **on the fly**

# Doc2vec

Distributed Memory Model of Paragraph Vectors (PV-DM)

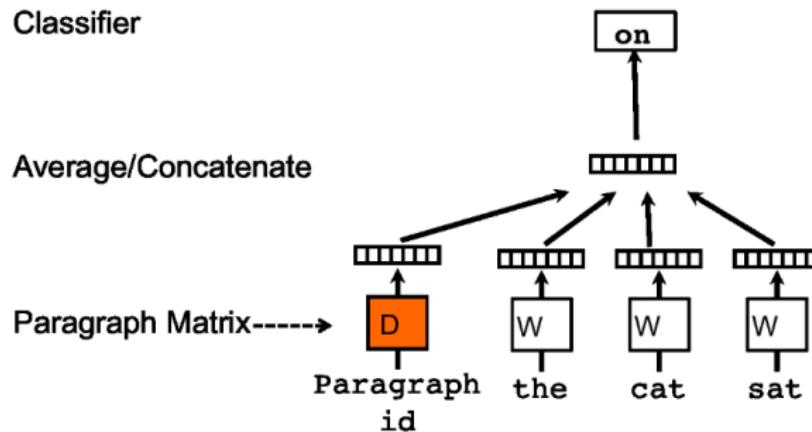
Derived from CBOW



# Doc2vec

Distributed Memory Model of Paragraph Vectors (PV-DM)

Derived from CBOW

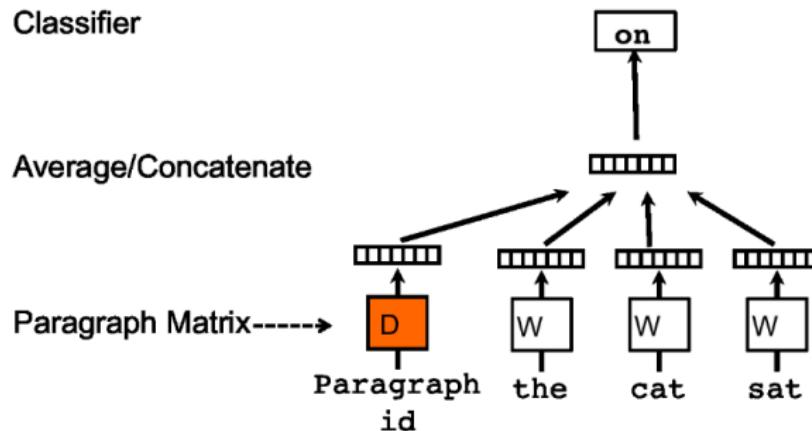


- Each column in the paragraph matrix is a vector representing one paragraph

# Doc2vec

Distributed Memory Model of Paragraph Vectors (PV-DM)

Derived from CBOW

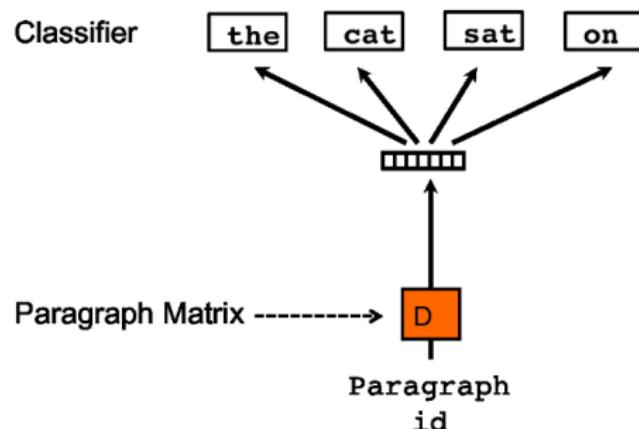


- Each column in the paragraph matrix is a vector representing one paragraph
- We can average or concatenate the word and paragraph vectors

# Doc2vec

Distributed Bag of Words version of Paragraph Vector (PV-DBOW)

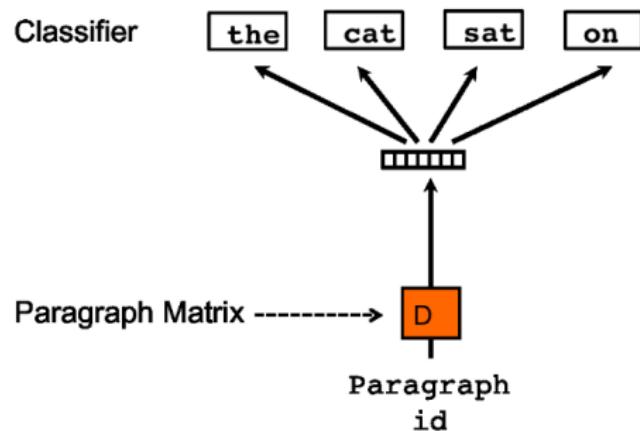
Similar to skip-gram



# Doc2vec

Distributed Bag of Words version of Paragraph Vector (PV-DBOW)

Similar to skip-gram

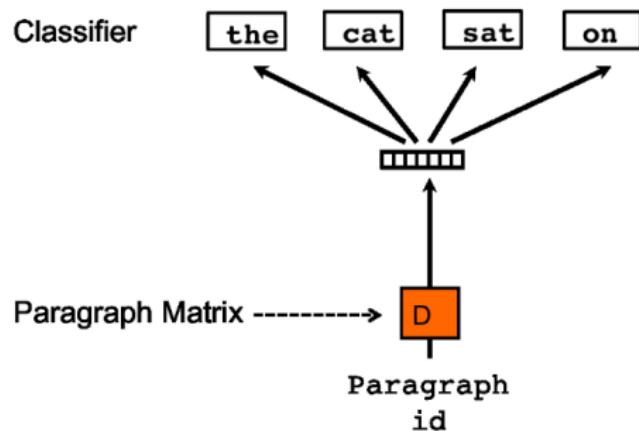


- Iteration: a text window and a random word from the text window are sampled, forming a classification task given the paragraph vector.
- No word vectors: faster + lower memory requirements

# Doc2vec

Distributed Bag of Words version of Paragraph Vector (PV-DBOW)

Similar to skip-gram



- Iteration: a text window and a random word from the text window are sampled, forming a classification task given the paragraph vector.
- No word vectors: faster + lower memory requirements

Let us see

## **Prologue to CNN and RNN**

# Prologue

- We have learned to build embedding spaces for words and texts

# Prologue

- We have learned to build embedding spaces for words and texts
- We are considering the neighborhood of the words (the bag)

# Prologue

- We have learned to build embedding spaces for words and texts
- We are considering the neighborhood of the words (the bag)
- We are not considering actual connections yet

# Prologue

- We have learned to build embedding spaces for words and texts
- We are considering the neighborhood of the words (the bag)
- We are not considering actual connections yet
- The downstream application is usually classification or regression

# Prologue

- We have learned to build embedding spaces for words and texts
- We are considering the neighborhood of the words (the bag)
- We are not considering actual connections yet
- The downstream application is usually classification or regression

**We will start heading towards text generation**

# Words have relations and influence each other

## Word order

$s_1$  = The dog chased the cat.

$s_2$  = The cat chased the dog.

# Words have relations and influence each other

Word order

$s_1$  = The dog chased the cat.

$s_2$  = The cat chased the dog.

$$\text{sim}(\text{tfidf}(s_1), \text{tfidf}(s_2)) = 1$$

# Words have relations and influence each other

Word order

$s_1$  = The dog chased the cat.

$s_2$  = The cat chased the dog.

$$\text{sim}(\text{tfidf}(s_1), \text{tfidf}(s_2)) = 1$$

$$\text{sim}(\text{wv}(s_1), \text{wv}(s_2)) = 1$$

# Words have relations and influence each other

Word order

$s_1$  = The dog chased the cat.

$s_2$  = The cat chased the dog.

$$\text{sim}(\text{tfidf}(s_1), \text{tfidf}(s_2)) = 1$$

$$\text{sim}(\text{wv}(s_1), \text{wv}(s_2)) = 1$$

$$\text{sim}(\text{dv}(s_1), \text{dv}(s_2)) = 1$$

# Words have relations and influence each other

Word order

$s_1$  = The dog chased the cat.

$s_2$  = The cat chased the dog.

$$\text{sim}(\text{tfidf}(s_1), \text{tfidf}(s_2)) = 1$$

$$\text{sim}(\text{wv}(s_1), \text{wv}(s_2)) = 1$$

$$\text{sim}(\text{dv}(s_1), \text{dv}(s_2)) = 1$$

But  $s_1$  and  $s_2$  are not the same!

# Words have relations and influence each other

Word order

$s_1$  = The dog chased the cat.

$s_2$  = The cat chased the dog.

$$\text{sim}(\text{tfidf}(s_1), \text{tfidf}(s_2)) = 1$$

$$\text{sim}(\text{wv}(s_1), \text{wv}(s_2)) = 1$$

$$\text{sim}(\text{dv}(s_1), \text{dv}(s_2)) = 1$$

But  $s_1$  and  $s_2$  are not the same!

## Word proximity

$s$  = His mother, besides her son's willingness to amend the issue,  
decided to punish him

# Words have relations and influence each other

Word order

$s_1$  = The dog chased the cat.

$s_2$  = The cat chased the dog.

$$\text{sim}(\text{tfidf}(s_1), \text{tfidf}(s_2)) = 1$$

$$\text{sim}(\text{wv}(s_1), \text{wv}(s_2)) = 1$$

$$\text{sim}(\text{dv}(s_1), \text{dv}(s_2)) = 1$$

But  $s_1$  and  $s_2$  are not the same!

## Word proximity

$s$  = His mother, besides her son's willingness to amend the issue,  
decided to punish him

mother... decided | son... him

---

(Lane et al., 2019, p. 220)

# Words have relations and influence each other

## Spatial relation

Consider the position of words  
(~written)

---

(Lane et al., 2019, p. 220)

# Words have relations and influence each other

## Spatial relation

Consider the position of words  
(~written)

## Temporal relation

Consider words as time series  
(~spoken)

# Words have relations and influence each other

## Spatial relation

Consider the position of words  
(~written)

## Temporal relation

Consider words as time series  
(~spoken)

→ fixed-width window

# Words have relations and influence each other

## Spatial relation

Consider the position of words  
(~written)

## Temporal relation

Consider words as time series  
(~spoken)

→ fixed-width window

**convolutional neural networks**

# Words have relations and influence each other

## Spatial relation

Consider the position of words  
(~written)

→ fixed-width window

## convolutional neural networks

## Temporal relation

Consider words as time series  
(~spoken)

→ ongoing (unk) amount of time

# Words have relations and influence each other

## Spatial relation

Consider the position of words  
(~written)

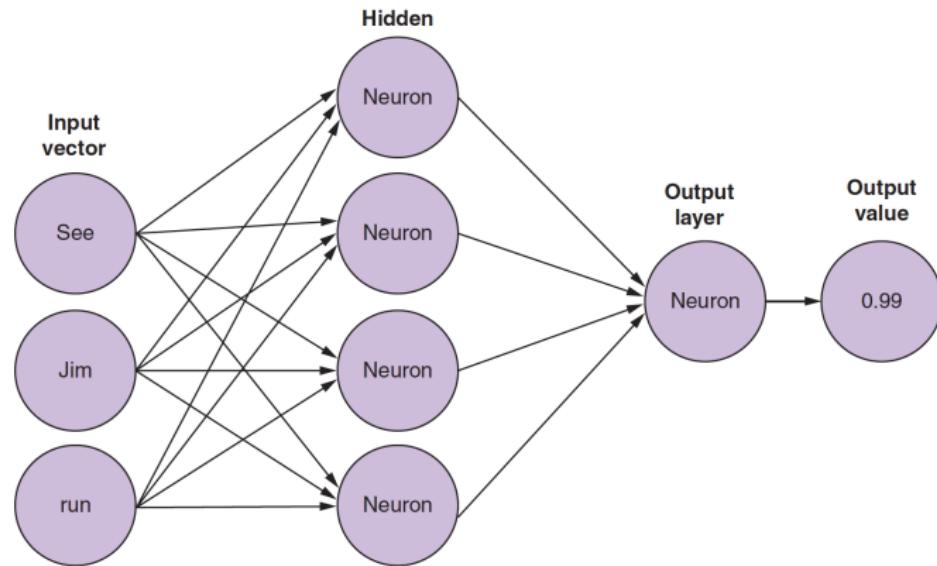
→ fixed-width window  
**convolutional neural networks**

## Temporal relation

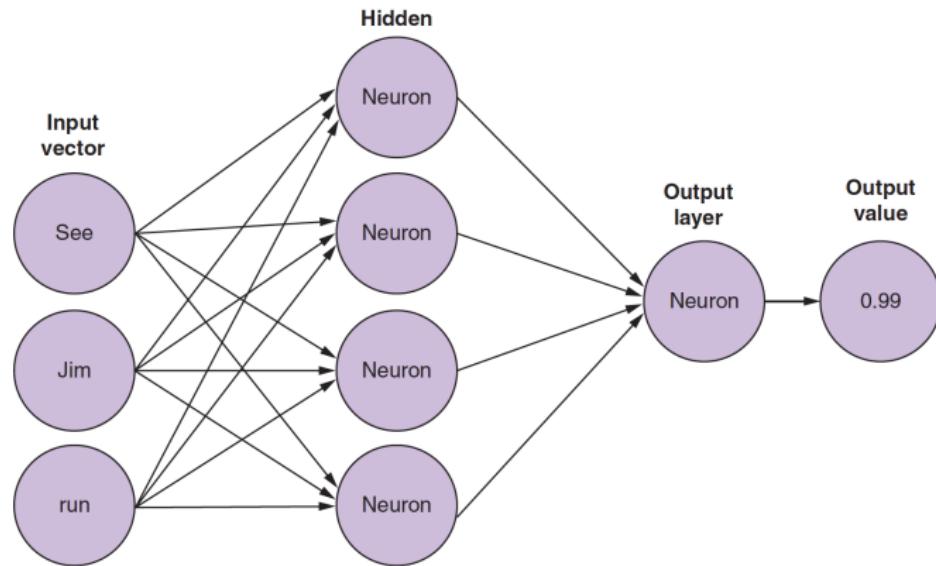
Consider words as time series  
(~spoken)

→ ongoing (unk) amount of time  
**recurrent neural networks**

# Multiple Input Words

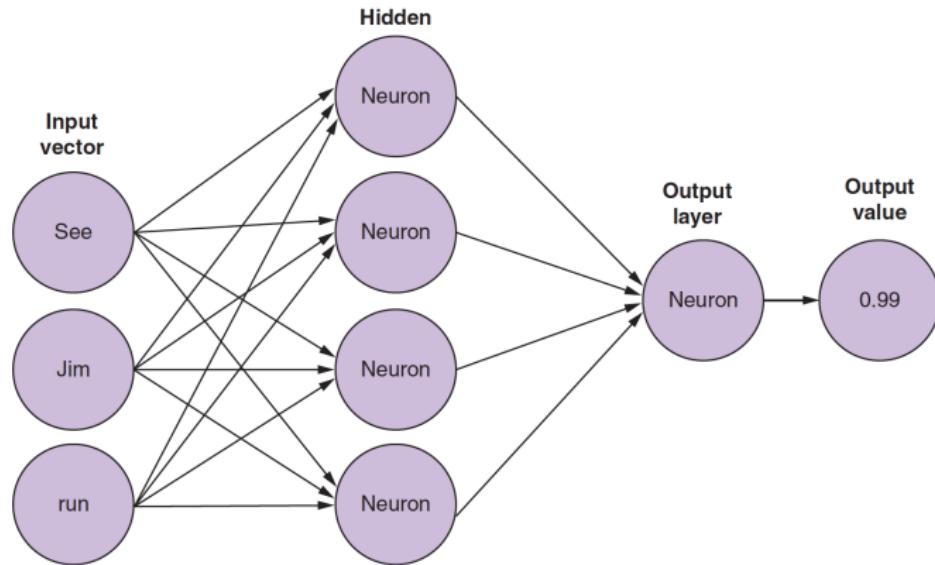


# Multiple Input Words



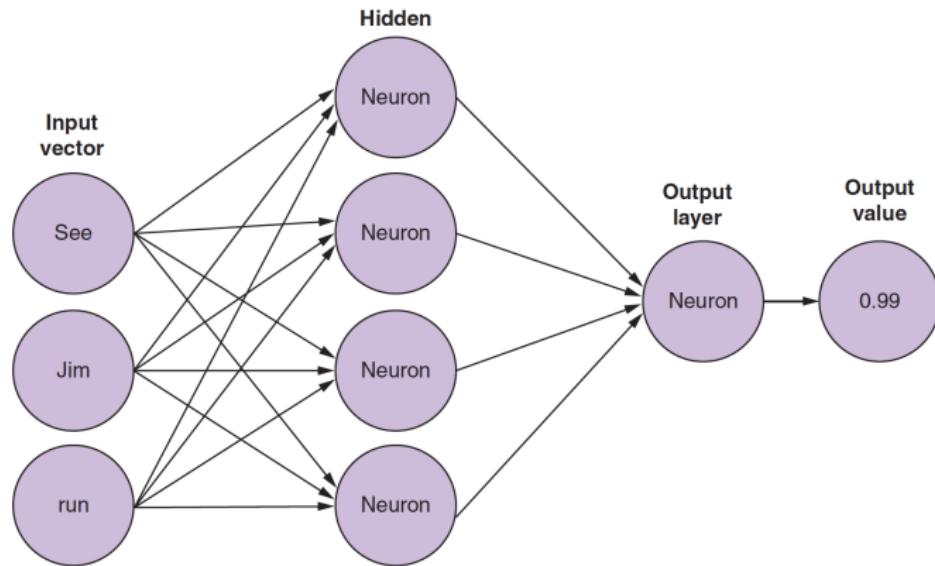
- Three tokens are passed at a time

# Multiple Input Words



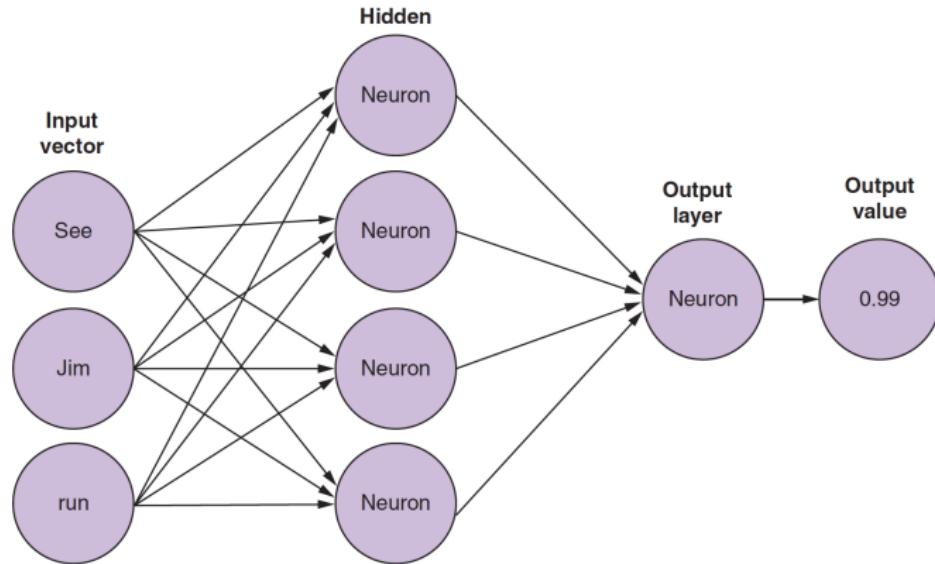
- Three tokens are passed at a time
- Two input alternatives
  - ▶ one-hot vector

# Multiple Input Words



- Three tokens are passed at a time
- Two input alternatives
  - ▶ one-hot vector
  - ▶ pre-trained word vector

# Multiple Input Words



- Three tokens are passed at a time
- Two input alternatives
  - ▶ one-hot vector
  - ▶ pre-trained word vector

See Jim run ≠ run See Jim (!)

# Back to Keras

## Sequential()

- Python class
- Neural network abstraction
- Grants access to the basic Keras API

# Back to Keras

## **Sequential()**

- Python class
- Neural network abstraction
- Grants access to the basic Keras API

## **Sequential.compile()**

- Builds the underlying weights
- Builds the and the interconnected relationships

# Back to Keras

## **Sequential()**

- Python class
- Neural network abstraction
- Grants access to the basic Keras API

## **Sequential.compile()**

- Builds the underlying weights
- Builds the and the interconnected relationships

## **Sequential.fit()**

- Computes the training errors (loss)
- Applies backpropagation (weight adjustment)

# Back to Keras

## Sequential()

- Python class
- Neural network abstraction
- Grants access to the basic Keras API

## Sequential.compile()

- Builds the underlying weights
- Builds the and the interconnected relationships

## Sequential.fit()

- Computes the training errors (loss)
- Applies backpropagation (weight adjustment)

## Some “cooking” hyperparameters

epochs number of iterations over the data

batch\_size number of instances before adjusting

optimizer function

# CNN

# Convolutional Neural Networks

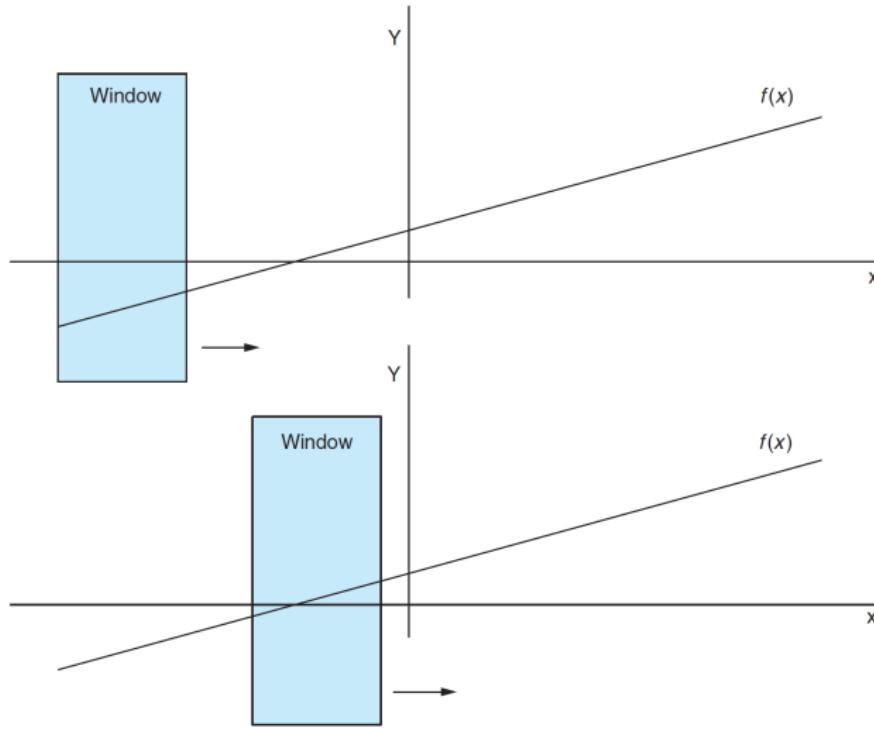
Sliding —or convolving<sup>1</sup>— a window over the sample

---

<sup>1</sup>To roll or wind together (Webster's)  
(Lane et al., 2019, p. 222)

# Convolutional Neural Networks

Sliding —or convolving<sup>1</sup>— a window over the sample



<sup>1</sup>To roll or wind together (Webster's)  
(Lane et al., 2019, p. 222)

# Convolutional Neural Networks

Back to the roots: image recognition

- Input: pixels of an image
- Output: the image contains x

---

[https:](https://blogs.nvidia.com/wp-content/uploads/2019/04/ADAS-IMG_0052.jpg)

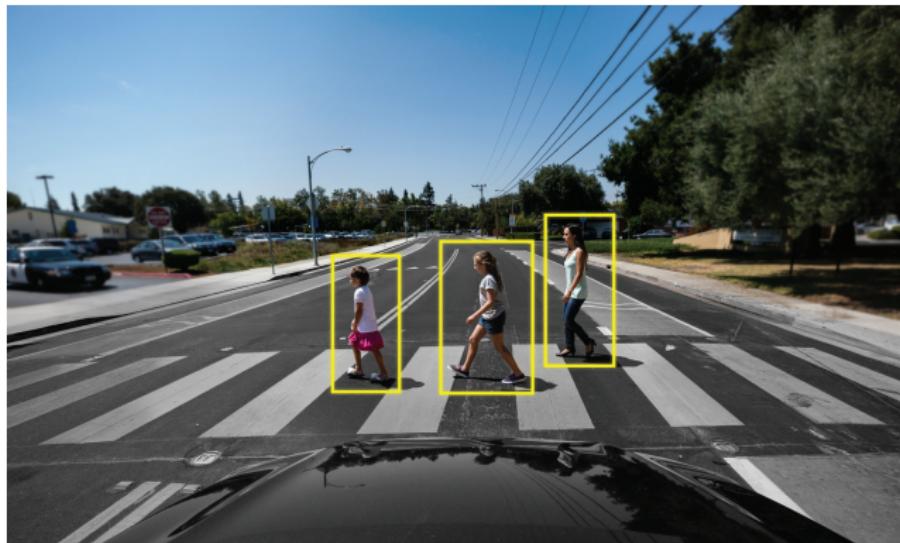
//blogs.nvidia.com/wp-content/uploads/2019/04/ADAS-IMG\_0052.jpg



# Convolutional Neural Networks

Back to the roots: image recognition

- Input: pixels of an image
- Output: the image contains x



---

[https:](https://blogs.nvidia.com/wp-content/uploads/2019/04/ADAS-IMG_0052.jpg)

//blogs.nvidia.com/wp-content/uploads/2019/04/ADAS-IMG\_0052.jpg



# Convolutional Neural Networks

When the input is an image

- B&W: [0,1] (with a smooth binariser)
- Grayscale: [0, 255]
- Colour: R: [0, 255] G: [0, 255] B: [0, 255]

# Convolutional Neural Networks

When the input is an image

- B&W: [0,1] (with a smooth binariser)
- Grayscaled: [0, 255]
- Colour: R: [0, 255] G: [0, 255] B: [0, 255]



---

(Lane et al., 2019, p. 223)

# Convolutional Neural Networks

When the input is an image

- B&W: [0,1] (with a smooth binariser)
- Grayscale: [0, 255]
- Colour: R: [0, 255] G: [0, 255] B: [0, 255]



0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27				
0	120	119	118	108	103	111	113	115	111	117	120	120	120	119	121	114	118	120	120	120	121	121	115	100	120	118	117				
1	111	109	111	106	107	118	120	110	106	117	120	119	119	119	121	114	118	119	118	119	118	114	119	109	102	122	114	117			
2	101	102	114	117	108	116	111	118	117	118	117	119	119	118	116	115	116	117	116	117	116	117	116	115	103	100	117	119	119		
3	109	102	114	117	108	116	111	118	117	118	117	119	119	118	116	115	116	117	116	117	116	117	116	115	103	100	117	119	119		
4	108	110	100	116	111	95	104	110	114	117	117	117	117	117	117	106	107	108	94	90	42	54	70	62	86	70	96	118			
5	103	112	109	98	100	93	111	112	115	113	117	117	117	117	116	116	112	118	116	114	110	43	81	110	102	119	114	119	117		
6	111	112	109	107	106	103	110	111	112	113	114	115	116	117	118	119	116	117	118	119	116	117	118	115	103	102	115	114	115		
7	111	111	109	105	103	110	120	103	106	111	115	115	116	110	107	103	113	114	115	115	48	87	105	105	114	111	116	115			
8	112	105	108	94	89	102	95	103	111	111	113	111	108	106	109	112	109	111	113	114	115	45	83	104	108	111	112	114	114		
9	112	106	108	86	95	109	104	103	108	104	108	109	114	114	113	108	112	113	113	114	43	80	104	111	109	113	114	113			
10	112	104	103	102	103	102	103	102	103	102	103	102	103	102	103	102	103	102	103	102	103	102	103	102	103	102	103	102			
11	110	106	93	96	108	107	110	109	111	112	108	107	111	112	111	111	107	111	111	112	110	38	78	109	109	108	111	114	103		
12	101	93	76	101	103	107	107	108	110	107	103	111	111	110	109	106	112	111	111	109	37	82	108	107	111	113	111	100			
13	98	92	99	115	108	108	111	106	100	99	106	108	109	110	107	106	109	108	109	107	37	78	106	103	106	108	103	98			
14	109	97	105	108	98	100	97	102	100	100	100	100	100	100	100	100	100	100	100	100	100	100	100	100	100	100	100	107			
15	84	69	92	97	85	95	92	99	95	98	100	107	107	107	108	106	104	108	107	105	29	74	107	107	110	106	99	109			
16	71	82	87	85	78	89	106	104	99	106	106	105	106	105	104	103	106	106	107	103	21	72	106	106	109	100	102	109			
17	67	87	64	68	84	89	98	96	99	104	104	104	105	103	101	105	103	103	106	102	23	76	103	103	106	98	108	101			
18	64	84	84	84	81	81	84	81	81	81	81	81	81	81	81	81	81	81	81	81	81	81	81	81	81	81	81	73			
19	60	71	77	77	78	80	88	92	91	93	96	96	101	100	101	100	98	101	101	104	92	13	64	93	89	81	89	81	79		
20	84	98	87	94	101	100	101	101	103	101	101	103	98	98	100	96	97	87	12	71	100	97	93	105	91	101					
21	92	80	88	99	100	98	100	100	97	98	97	97	96	92	91	92	93	96	87	13	79	105	95	98	96	89	100				
22	77	87	87	87	87	87	87	87	87	87	87	87	87	87	87	87	87	87	87	87	87	87	87	87	87	87	87	104			
23	81	87	83	84	89	89	91	87	90	92	93	95	94	96	94	94	95	98	94	93	84	7	74	89	86	90	70	81	73		
24	60	66	82	92	90	90	87	90	94	94	94	94	93	94	92	93	94	95	81	0	76	92	90	81	77	65	58				
25	87	81	83	86	87	84	84	90	92	92	92	92	92	92	91	92	91	92	77	4	73	91	92	81	95	96	95				
26	87	88	88	83	85	91	91	93	90	91	92	93	91	91	93	90	90	89	73	8	66	92	83	84	92	91	91				
27	81	86	88	89	85	91	91	89	89	88	88	89	89	84	83	79	80	87	84	0	89	77	90	91	90	98					

# Convolutional Neural Networks

When the input is an image

**An image is just a bunch of numbers**

# Convolutional Neural Networks

When the input is an image

## An image is just a bunch of numbers

- Appropriate as input for an NN

# Convolutional Neural Networks

When the input is an image

## An image is just a bunch of numbers

- Appropriate as input for an NN
- But one single pixel has no real meaning

# Convolutional Neural Networks

When the input is an image

## An image is just a bunch of numbers

- Appropriate as input for an NN
- But one single pixel has no real meaning

→ Sliding over fragments of the image

# Convolutional Neural Networks

When the input is an image

## An image is just a bunch of numbers

- Appropriate as input for an NN
- But one single pixel has no real meaning

## → Sliding over fragments of the image

The convolution defines a set of filters (aka kernels) to do just that

# Convolutional Neural Networks

When the input is an image

## An image is just a bunch of numbers

- Appropriate as input for an NN
- But one single pixel has no real meaning

## → Sliding over fragments of the image

The convolution defines a set of filters (aka kernels) to do just that

- Take “snapshots” of different areas of the image
- Process them, one at a time

# Convolutional Neural Networks

## Strides and filters

### Stride

- The distance “traveled” when sliding
- Yet another parameter
- Never bigger than the size of the filter → overlapping areas

# Convolutional Neural Networks

Strides and filters

## Stride

- The distance “traveled” when sliding
- Yet another parameter
- Never bigger than the size of the filter → overlapping areas

Sounds familiar?

# Convolutional Neural Networks

Strides and filters

## Stride

- The distance “traveled” when sliding
- Yet another parameter
- Never bigger than the size of the filter → overlapping areas

Sounds familiar? *n-grams!*

# Convolutional Neural Networks

Strides and filters

## Stride

- The distance “traveled” when sliding
- Yet another parameter
- Never bigger than the size of the filter → overlapping areas

Sounds familiar? *n-grams!*

## Filter

- $n \times m$  surfaces

# Convolutional Neural Networks

Strides and filters

## Stride

- The distance “traveled” when sliding
- Yet another parameter
- Never bigger than the size of the filter → overlapping areas

Sounds familiar? *n-grams!*

## Filter

- $n \times m$  surfaces
- Typically  $n = m = 3$  (often  $n \neq m$ )

# Convolutional Neural Networks

## Strides and filters

### Stride

- The distance “traveled” when sliding
- Yet another parameter
- Never bigger than the size of the filter → overlapping areas

Sounds familiar? *n-grams!*

### Filter

- $n \times m$  surfaces
- Typically  $n = m = 3$  (often  $n \neq m$ )
- Includes a set of weights (fix for the whole image)

# Convolutional Neural Networks

## Strides and filters

### Stride

- The distance “traveled” when sliding
- Yet another parameter
- Never bigger than the size of the filter → overlapping areas

Sounds familiar? *n-grams!*

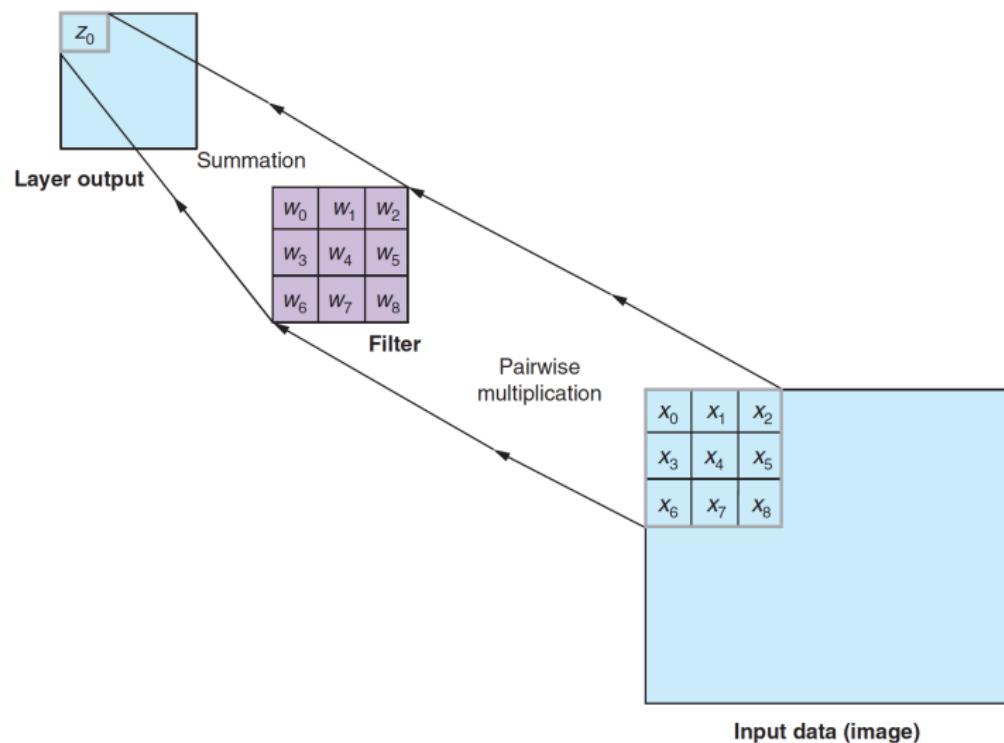
### Filter

- $n \times m$  surfaces
- Typically  $n = m = 3$  (often  $n \neq m$ )
- Includes a set of weights (fix for the whole image)
- Includes an activation function: usually ReLU

$$z = \max(\text{sum}(x * w), 0)$$

# Convolutional Neural Networks

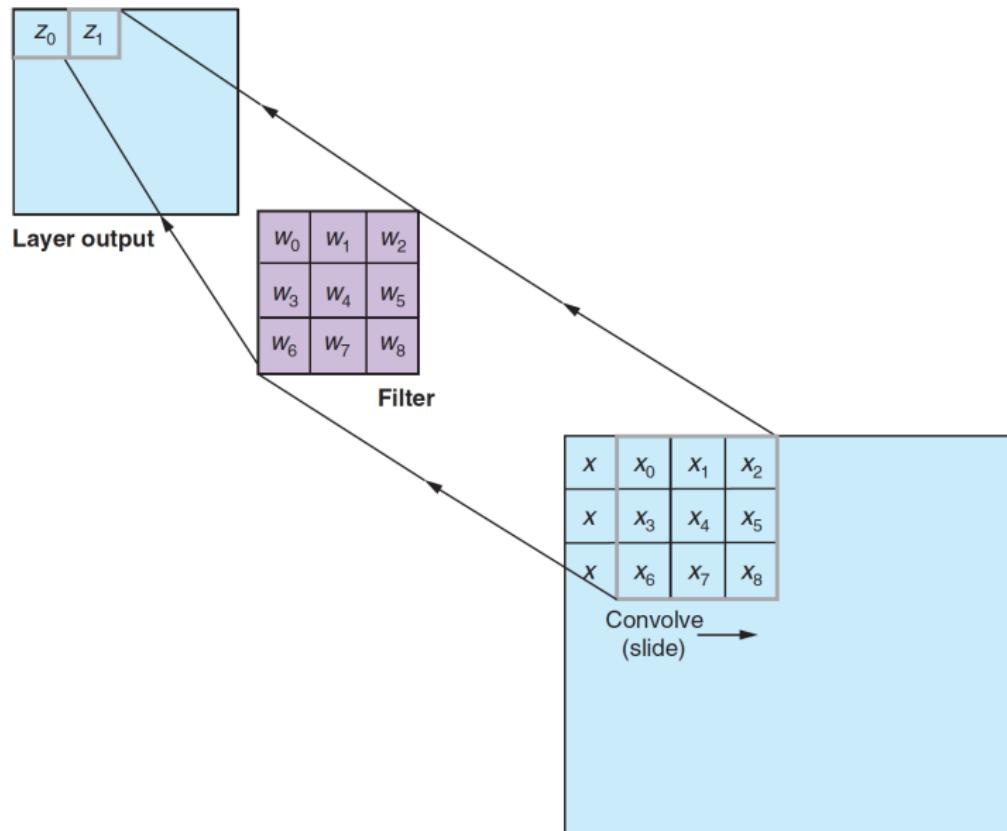
## Convolutional step



(Lane et al., 2019, p. 225)

# Convolutional Neural Networks

## Convolution



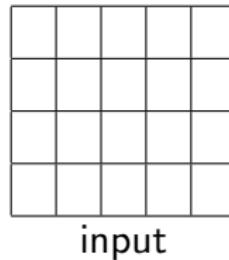
# Convolutional Neural Networks

Producing multiple images

- $k$  filters exist which carry out different operations
- Every filter will produce a new image, combination of source and filter

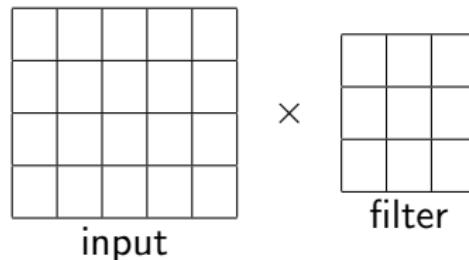
# Convolutional Neural Networks

## Padding



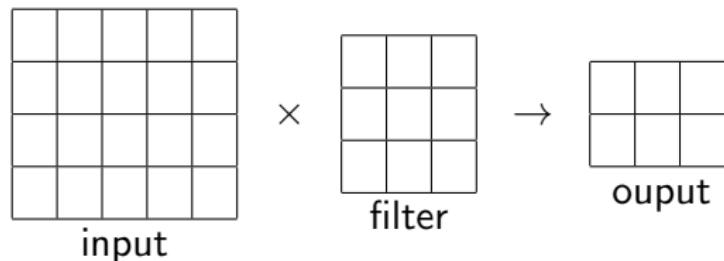
# Convolutional Neural Networks

## Padding



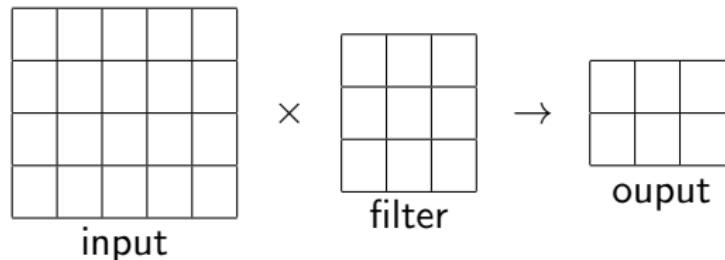
# Convolutional Neural Networks

## Padding



# Convolutional Neural Networks

## Padding

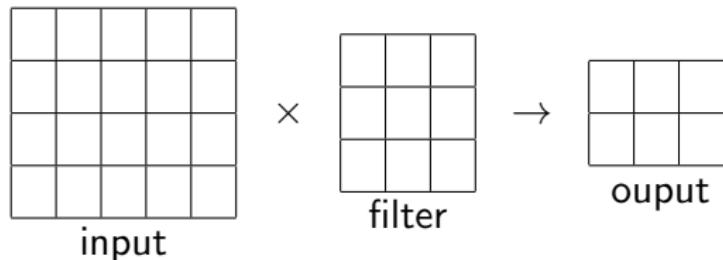


We are producing smaller images

"I don't care": Keras' argument padding='valid'

# Convolutional Neural Networks

## Padding



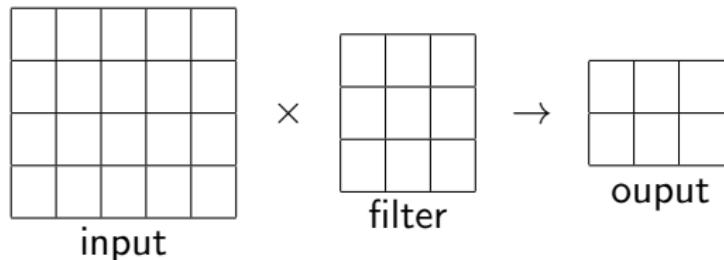
We are producing smaller images

"I don't care": Keras' argument padding='valid'

The edges of the image are undersampled

# Convolutional Neural Networks

## Padding



We are producing smaller images

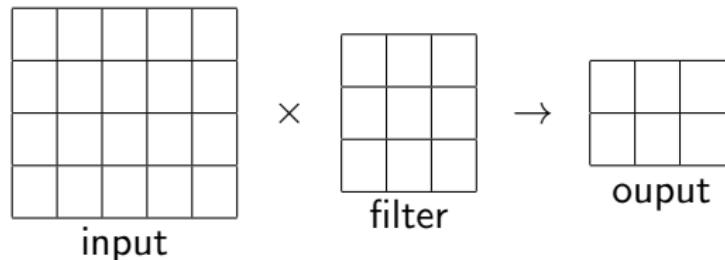
"I don't care": Keras' argument `padding='valid'`

The edges of the image are undersampled

"I do care": Keras' padding argument `padding='same'`

# Convolutional Neural Networks

## Padding



**We are producing smaller images**

“I don't care”: Keras' argument `padding='valid'`

The edges of the image are undersampled

“I do care”: Keras' padding argument `padding='same'`

In NLP **we care**

# Convolutional Neural Networks

## Pipeline

**Input:** an image, text

**Output:** a class, a real number

- Produce  $k$  new images through  $k$  filters

# Convolutional Neural Networks

Pipeline

**Input:** an image, text

**Output:** a class, a real number

- Produce  $k$  new images through  $k$  filters
- Wire the filtered images to a feed-forward

# Convolutional Neural Networks

## Pipeline

**Input:** an image, text

**Output:** a class, a real number

- Produce  $k$  new images through  $k$  filters
- Wire the filtered images to a feed-forward
- Proceed as usual

# Convolutional Neural Networks

Pipeline

**Input:** an image, text

**Output:** a class, a real number

- Produce  $k$  new images through  $k$  filters
- Wire the filtered images to a feed-forward
- Proceed as usual

**We can add multiple convolution layers**

A full path of learning layers and abstractions

# Convolutional Neural Networks

Pipeline

**Input:** an image, text

**Output:** a class, a real number

- Produce  $k$  new images through  $k$  filters
- Wire the filtered images to a feed-forward
- Proceed as usual

**We can add multiple convolution layers**

A full path of learning layers and abstractions

- Edges

# Convolutional Neural Networks

Pipeline

**Input:** an image, text

**Output:** a class, a real number

- Produce  $k$  new images through  $k$  filters
- Wire the filtered images to a feed-forward
- Proceed as usual

**We can add multiple convolution layers**

A full path of learning layers and abstractions

- Edges
- Shapes

# Convolutional Neural Networks

Pipeline

**Input:** an image, text

**Output:** a class, a real number

- Produce  $k$  new images through  $k$  filters
- Wire the filtered images to a feed-forward
- Proceed as usual

**We can add multiple convolution layers**

A full path of learning layers and abstractions

- Edges
- Shapes
- Colours

# Convolutional Neural Networks

Pipeline

**Input:** an image, text

**Output:** a class, a real number

- Produce  $k$  new images through  $k$  filters
- Wire the filtered images to a feed-forward
- Proceed as usual

**We can add multiple convolution layers**

A full path of learning layers and abstractions

- Edges
- Shapes
- Colours
- **Concepts**

# Convolutional Neural Networks

Pipeline

**Input:** an image, text

**Output:** a class, a real number

- Produce  $k$  new images through  $k$  filters
- Wire the filtered images to a feed-forward
- Proceed as usual

**We can add multiple convolution layers**

A full path of learning layers and abstractions

- Edges
- Shapes
- Colours
- **Concepts**

**What is learned**

- Good filters
- “Standard” weights

# Convolutional Neural Networks

Keras premier

```
from keras.models import Sequential  
from keras.layers import Conv1D  
  
model = Sequential()  
  
model.add(Conv1D(filters=16,  
                 kernel_size=3,  
                 padding='same',  
                 activation='relu',  
                 strides=1,  
                 input_shape=(100, 300))  
)
```

# Next time

- CNNs for NLP

# References

Lane, H., C. Howard, and H. Hapkem

2019. *Natural Language Processing in Action*. Shelter Island, NY:  
Manning Publication Co.

Le, Q. V. and T. Mikolov

2014. Distributed representations of sentences and documents.