

92586 Computational Linguistics

Lesson 2. Tokens

Alberto Barrón-Cedeño

Alma Mater Studiorum-Università di Bologna
a.barron@unibo.it @albarron_

04/10/2022



Table of Contents

Words

Normalisation

Representations

Words

Words

What is a word?

Speech The smallest sequence of phonemes that can be uttered in isolation with objective or practical meaning

Text Sequences of graphemes (“letters”) [...] delimited by spaces [...] or by other graphical conventions

<https://en.wikipedia.org/wiki/Word>

Simplistic definition

A word is a sequence of characters surrounded by spaces

Arguable, as multiple scholars claim; in particular across languages (Bender, 2013; Haspelmath, 2011)

Words

Lexicon

The set of all tokens (words!) in document d (or a corpus c)

Words

Tokenisers

We have a tokeniser, kindly provided by Church (1994)

```
tokens = re.findall('[A-Za-z]+', txt)
```

Python provides a “similar” tool

```
tokens = txt.split()
```

What if `txt` is the following?¹

```
txt = """Thomas Jefferson started building Monticello  
at the age of 26."""
```

`</>` **Let us see it working**

Words

Tokenisers

A better regular expression²

```
tokens = re.split(r'([-\\s.,;!?])+', txt)
```

What if we have the following text?

```
txt = "Monticello wasn't designated as UNESCO World Heritage Site  
until 1987"
```

`</>` **Let us see it working**

²Borrowed from Lane et al. (2019, p. 43)

Words

NLTK

- ▶ One of the leading platforms to work with human language data in python³
- ▶ Easy-to-use interfaces to over 50 corpora and lexical resources such as WordNet
- ▶ Suite of text processing libraries for classification, **tokenization**, stemming, tagging, parsing [...]

<http://www.nltk.org/>

³See spacy, stanza, huggingface

Words

Installing NLTK

```
$ pip install --user -U nltk
$ pip install --user -U numpy
$ python
>>> import nltk
```

Using (one of) the NLTK tokenisers

```
from nltk.tokenize import TreebankWordTokenizer
tokenizer = TreebankWordTokenizer()
sentence = "Monticello wasn't designated as UNESCO World
Heritage Site until 1987"
tokenizer.tokenize(sentence)
```

</> Let us see it working

Normalisation

Normalisation

Case folding

Ignoring differences in the spelling of a word which involves only capitalisation (Lane et al., 2019, p. 54)

```
# We know how to deal with this, don't we?
```

PROS Tea==tea; the vocabulary is smaller

CONS The Joker is not a character any longer

</> Let us see it working

Normalisation

Stemming

“Eliminate the small meaning differences of pluralisation or possessive endings of words or [...] verb form” (Lane et al., 2019, p. 57)

```
import re
def stem(phrase):
    return ' '.join([re.findall('^(.*ss|.*)?(s)?$',
                               word)[0][0].strip("'") for word in phrase.lower().split()])

stem('houses')
stem("Doctor House's calls")
stem("stress")
```

</> Let us see it working

Normalisation

Stemming: Porter and Snowball

Once again, people have developed (and released) more sophisticated stemming algorithms

<https://tartarus.org/martin/PorterStemmer/>

<http://snowball.tartarus.org/>

```
from nltk.stem.porter import PorterStemmer
stemmer = PorterStemmer()
' '.join([stemmer.stem(w).strip("'") for w in
"dish washer's washed dishes".split()])
```

Normalisation

Lemmatisation

Associating several words down to their semantic common root
(adapted from (Lane et al., 2019, p. 59))

PROS Stemming might alter the meaning of a word

CONS It is more expensive; it requires a knowledge base of synonyms and endings, and part-of-speech tags

Normalisation

Lemmatisation: re-use, re-use!

```
import nltk
nltk.download('wordnet')

from nltk.stem import WordNetLemmatizer
lemmatizer = WordNetLemmatizer()

lemmatizer.lemmatize("better")
lemmatizer.lemmatize("better", pos="a")
```


</> Let us see it working

Representations

Representations

Vectors

An (Euclidean) vector is an entity endowed with a magnitude (the length of the line segment (A, B)) and a direction (the direction from A to B).



[https://en.wikipedia.org/wiki/Vector_\(mathematics_and_physics\)](https://en.wikipedia.org/wiki/Vector_(mathematics_and_physics))
https://en.wikipedia.org/wiki/Vector_space

Representations

Bag of Words (BoW)

Turning words into numbers⁴

```
sentence = ""Thomas Jefferson began building  
Monticello at the age of 26.""
```

```
sentence_bow = {}  
for token in sentence.split():  
    sentence_bow[token] = 1  
sorted(sentence_bow.items())
```

⁴From (Lane et al., 2019, p. 35)

Representations

Bag of Words (BoW)

Using **pandas** (data structures for data analysis, time series, and statistics)⁵

```
import pandas as pd  
  
sentences = ""Thomas Jefferson began building Monticello at  
the age of 26.\n""  
sentences += ""Construction was done mostly by local masons  
and carpenters.\n""  
sentences += "He moved into the South Pavilion in 1770.\n"  
sentences += ""Turning Monticello into a neoclassical  
masterpiece was Jefferson's obsession.""  
corpus = {}  
for i, sent in enumerate(sentences.split('\n')):  
    corpus['sent{}'.format(i)] = dict((tok, 1) for tok in  
        sent.split())  
df = pd.DataFrame.from_records(corpus).fillna(0).astype(int).T  
df[df.columns[:10]]
```

</> Let us see it working

Representations

One-Hot Vectors

Turning words into numbers⁶

```
import numpy as np  
sentence = "Thomas Jefferson began building Monticello at  
the age of 26."  
token_sequence = str.split(sentence)  
vocab = sorted(set(token_sequence))  
, '.join(vocab)  
  
num_tokens = len(token_sequence)  
vocab_size = len(vocab)  
onehot_vectors = np.zeros((num_tokens, vocab_size), int)  
for i, word in enumerate(token_sequence):  
    onehot_vectors[i, vocab.index(word)] = 1  
  
, '.join(vocab)  
onehot_vectors
```

⁶From (Lane et al., 2019, p. 35)

Representations

One-Hot Vectors

Turning words into numbers⁷

```
import pandas as pd
pd.DataFrame(onehot_vectors, columns=vocab)
```

⁷From (Lane et al., 2019, p. 35)

References

- Bender, E. M.
2013. *Linguistic Fundamentals for Natural Language Processing: 100 Essentials from Morphology and Syntax*. Morgan & Claypool Publishers.
- Church, K.
1994. UNIX for poets.
- Haspelmath, M.
2011. The indeterminacy of word segmentation and the nature of morphology and syntax. *Folia Linguistica*, 45.
- Lane, H., C. Howard, and H. Hapkem
2019. *Natural Language Processing in Action*. Shelter Island, NY: Manning Publication Co.