

91258 Natural Language Processing

Lesson 17. Bidirectional RNN → Long Short-Term Memory Networks

Alberto Barrón-Cedeño

Alma Mater Studiorum-Università di Bologna
a.barron@unibo.it @albarron_

27/11/2023



Table of Contents

End of Chapter 8, Chapter 9 of Lane et al. (2019)

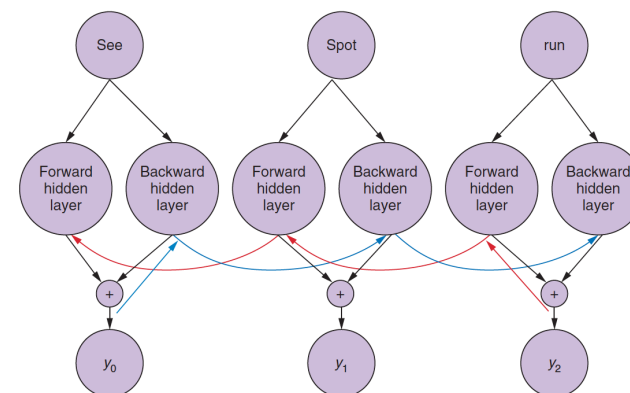
Left and right context

Not only the previous context is important to understand the *current* token

They wanted to pet the dog whose fur was brown.

- Descriptions and relevant information often come later
- A standard RNN neglects information from the *future*

Bidirectional recurrent neural network




- We *arrange* 2 RNNs:
 - one takes the input as usual
 - the other takes the backward input
 - + means concatenation

Implementation difference

```
# Adding one bidirectional recurrent layer

model.add(Bidirectional(SimpleRNN(
    num_neurons,
    return_sequences=True),
    input_shape=(maxlen, embedding_dims))
)
```

 Let us see

LSTMs

BiRNN zoom into results

Accuracies after 2 epochs

units	Acc	Acc _{val}
50	0.8156	0.7662
40	0.8244	0.7540
30	0.8259	0.7874
20	0.8072	0.8076
10	0.8007	0.8016
5	0.7973	0.8006
1	0.7070	0.7822

* remember we had used 50 units last time for the RNN

Short effect from the past

The effect of token x_t dilutes significantly as soon as in $t + 2$

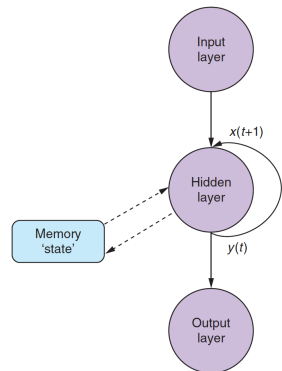
Consider the following —fairly plausible— texts. . .

The young woman went to the movies with her friends.

The young woman, having found a free ticket on the ground, went to the movies.

- ▶ In both cases, **went** is the main verb
- ▶ A (Bi)RNN would hardly reflect that in the second case
- ▶ We need an architecture able to "remember" the entire input

State: the memory of an LSTM



- ▶ The memory state contains attributes
- ▶ The attributes are updated with every instance
- ▶ The *rules* of the state are trained NNs


Now we have two learning objectives:

- ▶ Learn to predict the target labels
- ▶ Learn to identify what has to be *remembered*

(Lane et al., 2019, p. 276)


Unrolled LSTM

- ▶ Activation from $t - 1$ plus memory state
- ▶ The memory state sends a vector with the state of each **LSTM cell**, of cardinality number_of_units



(Lane et al., 2019, p. 277)


The LSTM cell (layer)



Input: $\text{output}_{t-1} \oplus \text{input}_t$

Gates: a FF layer + an activation function **each**

LSTM Forget Gate



Input:

$[x_{[t,0]}, x_{[t,1]}, \dots, x_{[t,299]}, h_{[t-1,0]}, h_{[t-1,1]}, \dots, h_{[t-1,49]}, 1]$

Forget: How much of the memory should be erased

—forgetting long-term dependencies as new ones arise

$351 * 50 = 17,550$ parameters


Feed-forward NN with sigmoid activation function:

$[0, 1]$

(Lane et al., 2019, p. 280)

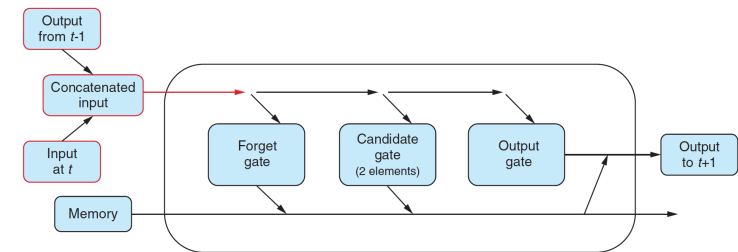
LSTM Forget Gate

Forget is a mask:



(Lane et al., 2019, p. 282)

LSTM Candidate Gate

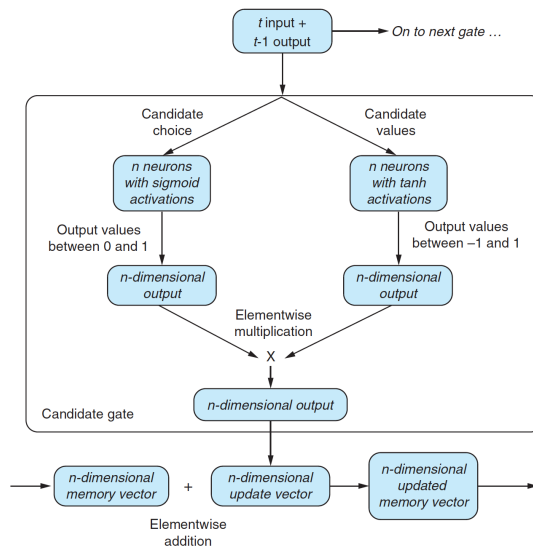


Input:

$[x_{[t,0]}, x_{[t,1]}, \dots, x_{[t,299]}, h_{[t-1,0]}, h_{[t-1,1]}, \dots, h_{[t-1,49]}, 1]$

Candidate: How much to augment the memory —what to remember and where to do it

LSTM Candidate Gate

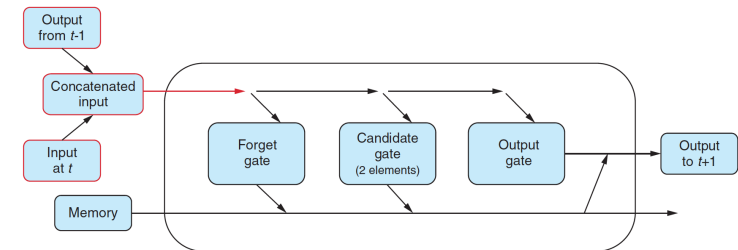


Candidate choice
Which values should be updated (\sim forget)

Candidate values
Computes those new values

(Lane et al., 2019, p. 283)

LSTM Output Gate



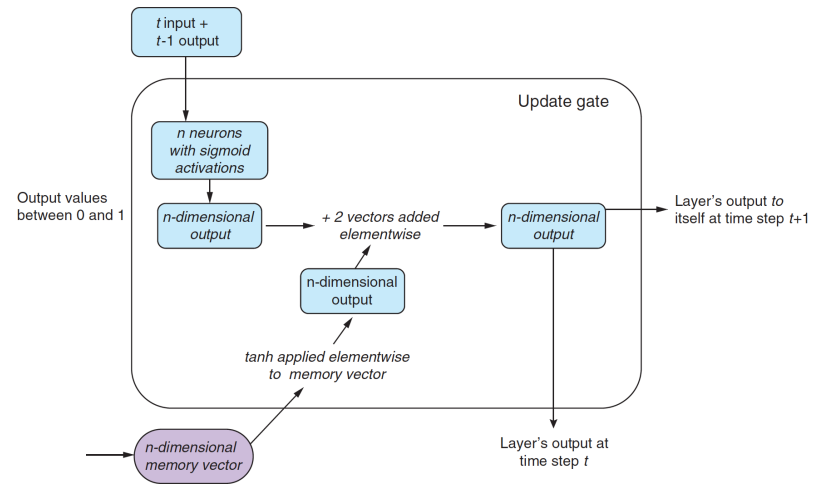
Input:

$[x_{[t,0]}, x_{[t,1]}, \dots, x_{[t,299]}, h_{[t-1,0]}, h_{[t-1,1]}, \dots, h_{[t-1,49]}, 1]$

Output: produces the output vector —both for the actual task and back to the memory

- sigmoid to the input
- tanh to the state

LSTM Output Gate



* The figure says "added". It is a product

(Lane et al., 2019, p. 284)

LSTM: Wrapping Up

- The *main* network uses the output of the memory in the same fashion as in a RNN
- The memory *decides* what to keep/feed to the network
- The weights of the memory are also learned by back-propagation

📖 Let us see

LSTM: Result

arch	units	Acc	Acc _{val}
BiRNN	50	0.8156	0.7662
BiRNN	40	0.8244	0.7540
BiRNN	30	0.8259	0.7874
BiRNN	20	0.8072	0.8076
BiRNN	10	0.8007	0.8016
BiRNN	5	0.7973	0.8006
BiRNN	1	0.7070	0.7822
LSTM	50	0.8692	0.8678

References

Lane, H., C. Howard, and H. Hapkem
2019. *Natural Language Processing in Action*. Shelter Island, NY: Manning Publication Co.