

91258 Natural Language Processing

Lesson 19. Into Transformers¹

Alberto Barrón-Cedeño

Alma Mater Studiorum-Università di Bologna
a.barron@unibo.it @albarron_

04/12/2023



¹Partially based on medium.com/inside-machine-learning/what-is-a-transformer-d07dd1fbec04

Table of Contents

Sequence to Sequence Models

Attention is all you need

Transformers

Sequence to Sequence Models

Seq2Seq


*Seq2Seq models **transform** a sequence of elements (e.g., the words in a sentence) into another sequence*

Examples of problems that fit Seq2Seq?

- ▶ Text simplification
- ▶ Paraphrasing
- ▶ Machine translation

Seq2Seq

Encoder-Decoder architecture



Encoder takes the input sequence and maps it into a higher-dimensional space (vector)

Decoder turns the vector into an output sequence (language, symbols, copy of the input²)

²Smaller vector for compression
(Lane et al., 2019, 315)

Seq2Seq

Intuition³

- ▶ I need to translate texts from Italian to English
- ▶ I have two *translators*: Alice and Bob
 - ▶ Alice speaks Italian, but not English
 - ▶ Bob Speaks English, but not Italian
 - ▶ Both speak (just a bit of!) Spanish

What do I need to get Alice and Bob to translate properly together?

I need to teach them better Spanish

Alice is my encoder

Spanish *is* the language of my thought vector

Bob is my decoder

I need to learn (train) the model to encode/decode the text

³From medium

Seq2Seq

Noisy Channel




Diagram from Jurafsky's <https://image1.slideserve.com/1844322/the-noisy-channel-model-for-mt-1.jpg>

Attention is all you need

Attention (Vaswani et al., 2017)

The attention-mechanism looks at an input sequence and decides, at each step, which **other parts** of the sequence are important⁴

Encoder (LSTM) uses the attention mechanism to take into account several other inputs for each element in the input

Decoder (LSTM) takes both the encoded sentence and the weights from the attention mechanism.

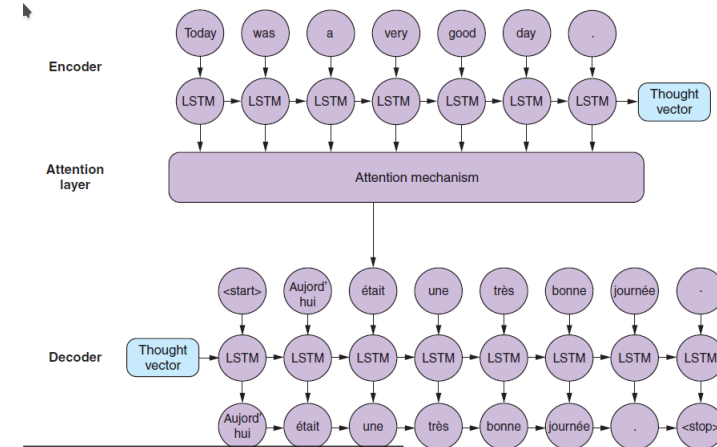
⁴Memory in an LSTM rings a bell?

Transformers

Attention

Sequence Labelling

- Part-of-speech tagging
- Dependency parsing
- Named entity recognition



(Lane et al., 2019, 334)

Transformers

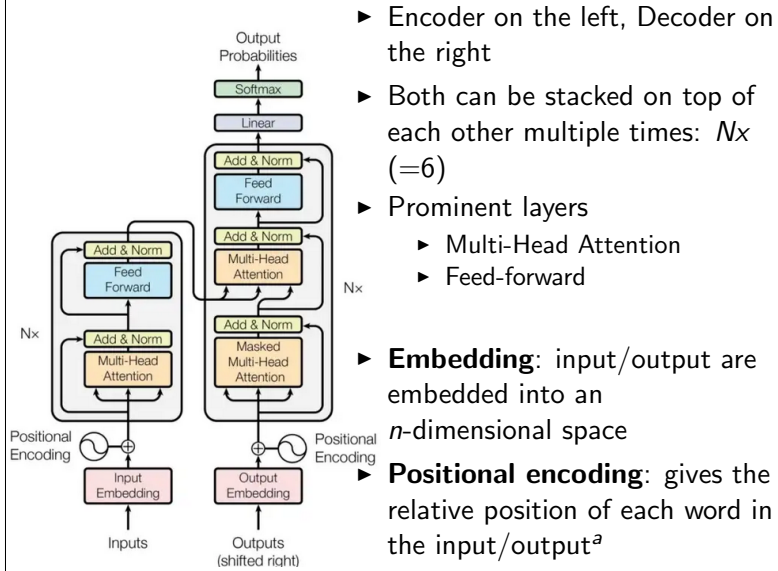
A Transformer [...] helps in transforming one sequence of input into another depending on the problem statement. Examples:

- Translation from one language to another
- Paraphrasing
- Question answering

No recurrent neural networks in this case

Transformers

Architecture (Vaswani et al., 2017)



- Encoder on the left, Decoder on the right
- Both can be stacked on top of each other multiple times: $N \times$ ($=6$)
- Prominent layers
 - Multi-Head Attention
 - Feed-forward
- **Embedding**: input/output are embedded into an n -dimensional space
- **Positional encoding**: gives the relative position of each word in the input/output^a

^aThis is not a recurrent network

Transformers

Attention

An attention function can be described as mapping a query and a set of key-value pairs to an output, where the query, keys, values, and output are all vectors. The output is computed as a weighted sum of the values, where the weight assigned to each value is computed by a compatibility function of the query with the corresponding key


(Vaswani et al., 2017)

Transformers

Muli-head attention (Vaswani et al., 2017)

$$Attention(Q, K, V) = softmax\left(\frac{Q \cdot K^T}{\sqrt{d_k}}\right) V$$

Scaled Dot-Product Attention




- Q queries: vector representation of one word in the sequence
- K keys: to the vector representations for all the words in the sequence
- V values of the vector representations for all the words in the sequence (same as Q)^a
- d_k Dimension of Q and K

$Attention(Q, K, C)$ weights on the values

^aThere is a trick here: actually, we have $q \in Q \forall q$ in the sequence

Transformers

Muli-head attention (Vaswani et al., 2017)

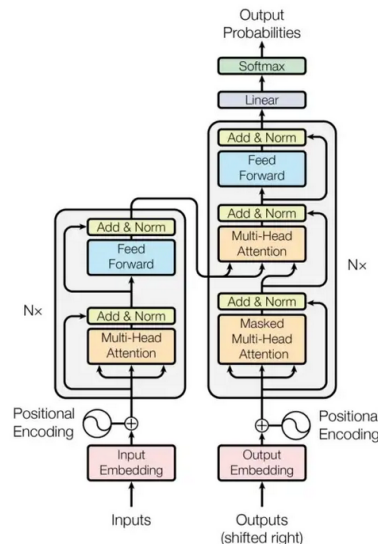


“Linearly project[ing] the queries, keys and values h times with different, learned linear projections to d_k , d_k and d_v dimensions

Matrices W that are learned (rings a bell?)

Transformers

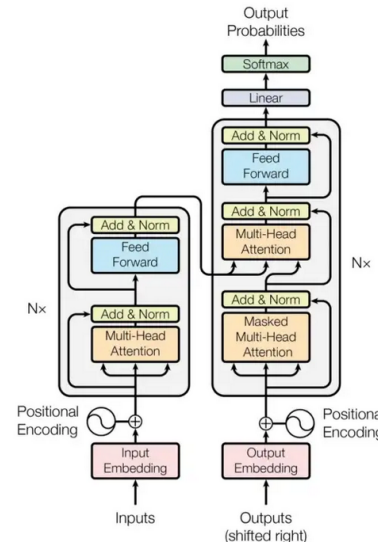
Attention in words



- The weights define how each word in sequence Q is influenced by all other words in the sequence (K)
- SoftMax distributes the weight over all words ($\sum_K = 1$)
- The weights are applied to all the words in sequence V
- Matrices Q , K , and V are different for each attention module
- The module connecting encoder and decoder takes into account the encoder input-sequence together with the decoder input-sequence up to a given position

Transformers

Training



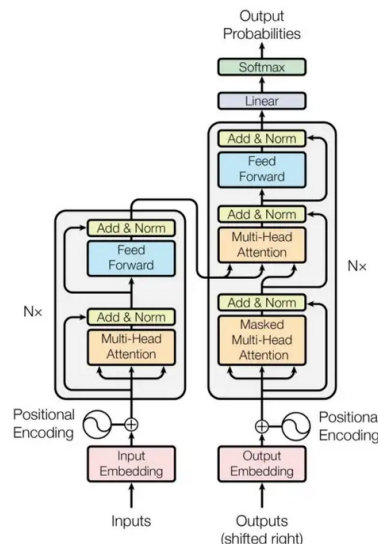
input_e $x_0 x_1 x_2 x_3 \dots x_{|X|}$
 input_d $y_1 y_2 y_3 y_4 \dots y_{|Y|+1}$

Why shifting input_d?

We want to learn that, given the encoder sequence and a particular decoder sequence (both seen already by the model), we have to predict the next word/character (otherwise, the model learns to copy the input_d)

Transformers

Inference



- Input the full input_e and an empty input_d (start-of-sentence token)
- Get the first element of the output produced
- Input the full input_e and start-of-sentence + first output element
- Repeat until end-of-sentence

References I

- Lane, H., C. Howard, and H. Hapkem
 2019. *Natural Language Processing in Action*. Shelter Island, NY: Manning Publication Co.
- Vaswani, A., N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin
 2017. Attention is all you need.