

92586 Computational Linguistics

9. “More than One” Neuron

Alberto Barrón-Cedeño

Alma Mater Studiorum-Università di Bologna
a.barron@unibo.it @albarron_

16/04/2020



Previously

- The perceptron
- Intro to neural networks

Table of Contents

1 Backpropagation (brief)

2 Keras

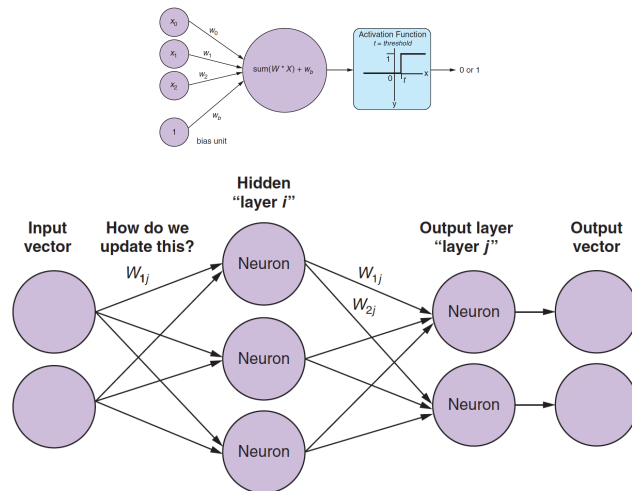
3 Some Guidelines

Chapter 5 of Lane et al. (2019)

Backpropagation (brief)

Weight Updating


Learning in a “simple” perceptron¹ vs a fully-connected network



(Lane et al., 2019, p. 158, 168)

¹Remember: aka linear regression

Backpropagation (of the errors)



- The error is computed on the output vector
- How much error did W_{1i} “contribute”?²
- “Path”: $W_{1i} \rightarrow [W_{1j}, W_{2j}] \rightarrow output$

²Notice that the first W_{1j} should be W_{1i}

Backpropagation (of the errors)


A better activation function

Step function: $f(\vec{x}) = \begin{cases} 1 & \text{if } \sum_{i=0}^n x_i w_i > threshold \\ 0 & \text{otherwise} \end{cases}$

Sigmoid function: non-linear and continuously differentiable

$$S(x) = \frac{1}{1 + e^{-x}} \quad (1)$$

Let us see



Non-linear \rightarrow model non-linear relationships

Continuously differentiable \rightarrow partial derivatives wrt various variables to update the weights

Backpropagation

Differentiating to adjust

Squared error (in (Lane et al., 2019, p. 171) they say this is MSE; wrong)

$$SE = (y - f(x))^2 \quad (2)$$

Mean squared error

$$MSE = \frac{1}{n} \sum_{i=1}^n (y - f(x))^2 \quad (3)$$

Calculus chain rule

$$f(g(x))' = F'(x) = f'(g(x)) g'(x) \quad (4)$$


With (4) we can find the derivative of the actfunc \forall neuron wrt its input.

Plain words: find the contribution of a weight to the error and adjust it!


(no further math)

Backpropagation (of the errors)

~Gradient descent: minimising the error



Convex error curve



Non-convex error curve

(Lane et al., 2019, p. 173–174)

Addressing Local minima

Batch learning

- Aggregate the error for the batch
- Update the weight at the end
- → hard to find global minimum

Stochastic gradient descent

- Look at the error for each single instance
- Update the weights right away
- → more likely to make it to the global minimum

Mini-batch

- Much smaller batch, combining the best of the two worlds
- → Fast as batch, resilient as stochastic gradient descent

Important parameter: learning rate α

A parameter to define at what extent should we “correct” the error

Keras

Some Available Libraries

There are many high- and low-level libraries in many languages

- **PyTorch**
Community-driven; <https://pytorch.org/>
- **Theano**
MILA (UdeM); www.deeplearning.net/software/theano/³
- **TensorFlow**
Google Brain; <https://www.tensorflow.org/>
- **Others**

We will use **Keras**; <https://keras.io/>

³Non active


What is Keras

- High-level wrapper with an accessible API for Python
- Gives access to three alternative backends
 - ▶ Theano
 - ▶ TensorFlow
 - ▶ CNTK (MS)

Keras


Logical exclusive OR in Keras

input		output
0	0	0
0	1	1
1	0	1
1	1	0

 Let us see

- First dense layer
 - ▶ 2 inputs, 10 neurons
 - ▶ 30 parameters
 - ▶ $2 \times 10 \rightarrow 20$
 - ▶ But we also have the bias! That's 10 more weights
- Second dense layer
 - ▶ 10 inputs, 1 neuron
 - ▶ 11 parameters

Now we can compile the model

 Let us see

Some Guidelines

“Design Decisions”

Activation functions

- Sigmoid
- ReLU (rectified linear unit)
- tanh (hyperbolic tangent)

Learning rate

- Choosing one in advance
- Use alertmomentum to perform dynamic adjustments

Dropout

- Ignore randomly-chosen weights in a training pass to prevent overfitting

Regularisation

- Dampen a weight from growing/shrinking too far from the rest to prevent overfitting

Normalisation

Example House classification.

Input number of bedrooms, last selling price

Output Likelihood of selling

Vector `input_vec` = [2, 90000]

All input dimensions should have comparable values

Ideally, all features should be in the range $[-1, 1]$ or $[0, 1]$

Typical normalisation: mean normalization, feature scaling, coefficient of variation

NLP uses TF-IDF, one-hot encoding, word2vec (already normalised!)

References

Lane, H., C. Howard, and H. Hapkem

2019. *Natural Language Processing in Action*. Shelter Island, NY: Manning Publication Co.