

92586 Computational Linguistics

10. Word2vec

Alberto Barrón-Cedeño

Alma Mater Studiorum-Università di Bologna
a.barron@unibo.it @albarron_

21/04/2020



Previously

- Introduction to Neural Networks
- First Keras neural network
- Considerations when building/training a network

Table of Contents

- 1 Introduction
- 2 Word Vectors
- 3 Computing word2vec representations

Chapter 6 of Lane et al. (2019)

Introduction

Introduction

Previously

- Each token represents one dimension (BoW)
- Document- and corpus-based statistics (TF-IDF)
- Dimensional reduction (LSA)

Drawbacks

- Ignoring the (nearby) context of a word
- Ignoring the overall meaning of a statement

Introduction

Word vectors. Numerical vector representations of word semantics, or meaning, including literal and implied meaning (Lane et al., 2019, p. 182)

Math with words

q = “She invented something to do with physics in Europe in the early 20th century”

```
answer_vector = wv['woman'] + wv['Europe'] + \
                wv['physics'] + wv['scientist']
```

Even better:

```
answer_vector = wv['woman'] + wv['Europe'] + \
                wv['physics'] + wv['scientist'] - \
                wv['male'] - wv['man']
```

Word Vectors

Word Vectors

Intuition

Word2vec Mikolov et al. (2013)

- Learns the meaning of words by processing a large corpus¹
- The corpus is not labeled
- It is **unsupervised**

Can we train a neural network to predict word occurrences near a target w ?

We don't care about the prediction (that's nice, but not important right here). We care about the resulting **internal representation**

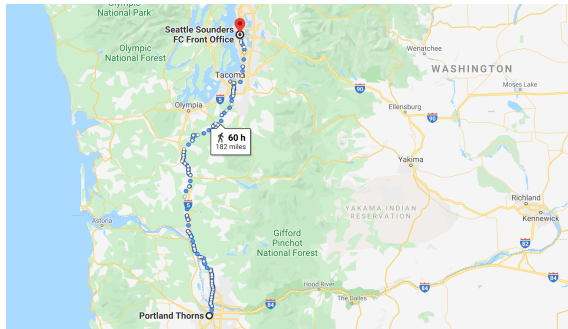
¹And I mean large; e.g., 100B words from Google News Groups

Word Vectors

Vector Algebra (again)

Portland Timbers + Seattle – Portland =?

$output_vector = wv['Seattle'] + wv['Portland Timbers'] - wv['Portland']$



Word Vectors

Vector Algebra (again)

- word2vec transforms token-occurrence vectors into lower-dimensional vectors
- The dimension is usually in the 100s (e.g., 100, 200)

Typical process

Input: Text

Output: Text

- 1 Compute vectors
- 2 Do algebra
- 3 Map back to text

Word Vectors

Some “typical” operations/properties

Gender king + woman – man → queen

Pl/Sg $\vec{x}_{coffee} - \vec{x}_{coffees} \approx \vec{x}_{cup} - \vec{x}_{cups} \approx \vec{x}_{cookie} - \vec{x}_{cookies}$

Locations San Francisco – California + Colorado → Denver

Culture tortellini – Bologna + Valencia → paella ?

Computing word2vec representations

Alternatives to Build word2vec representations

skip-gram

Input one (target) word

Output context words

CBOW (continuous bag-of-words)

Input context words

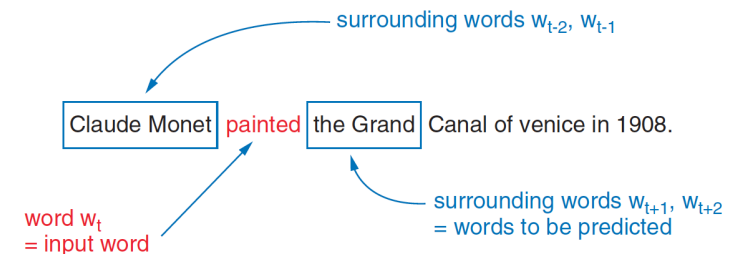
Output one target word

Skip-Gram

Definition Skip-grams are n -grams that contain gaps (skips over intervening tokens)

Input one word


Output context words



(Lane et al., 2019, p. 192)

Skip-Gram

Neural Network Structure



- n is the number of vector dimensions in the model
- M is the number of input/output neurons; $M = |\text{vocabulary}|$
- The output activation function is a **softmax** (typical in multi-class problems; $\sum_M = 1.0$)

Skip-Gram

Learning the Representations (1/3)

- Window size: 2 words \rightarrow 5-grams
- Input: each token, from left to right
- Output: the context on the left and right (one at a time)

$$S = w_1 w_2 w_3 w_4 w_5 w_6 w_7 w_8 w_9 w_{10}$$

$$[\dots] w_{t-2} w_{t-1} \underline{w_t} w_{t+1} w_{t+2} [\dots]$$

Skip-Gram

Learning the Representations (2/3)

Example: "Claude Monet painted the Grand Canal of Venice in 1908."

| input w_t | expected output | | | |
|----------------|-----------------|-----------|-----------|-----------|
| | w_{t-2} | w_{t-1} | w_{t+1} | w_{t+2} |
| Claude | | | Monet | painted |
| Monet | | Claude | painted | the |
| painted | Claude | Monet | the | Grand |
| the | Monet | painted | Grand | Canal |
| Grand | painted | the | Canal | of |
| Canal | the | Grand | of | Venice |
| of | Grand | Canal | Venice | in |
| Venice | Canal | of | in | 1908 |
| in | of | Venice | 1908 | |
| 1908 | Venice | in | | |

(Lane et al., 2019, p. 194)

Skip-Gram

Learning the Representations (3/3)

Training

- The input/output is a one-hot vector
- $n - 1$ iterations when using n -grams:

$[\dots] w_{t-2} w_{t-1} \underline{w_t} w_{t+1} w_{t+2} [\dots]$

| i | input | output | i | input | output | i | input | output |
|-----|-------|-----------|-----|-----------|-----------|-----|-----------|-----------|
| 0 | w_t | w_{t-2} | 4 | w_{t+1} | w_{t-1} | 8 | w_{t+2} | w_t |
| 1 | w_t | w_{t-1} | 5 | w_{t+1} | w_t | 9 | w_{t+2} | w_{t+1} |
| 2 | w_t | w_{t+1} | 6 | w_{t+1} | w_{t+2} | 10 | w_{t+2} | w_{t+3} |
| 3 | w_t | w_{t+2} | 7 | w_{t+1} | w_{t+3} | 11 | w_{t+2} | w_{t+4} |

- To simplify the loss calculation, the softmax is converted to one-hot

Skip-Gram

Outcome


- The output layer can be ignored²
- Semantically similar words have similar vectors—they were trained to **predict similar contexts**
- The weights from input to hidden layer are used to compute **embeddings**

$$wv_w = \text{dot}(\text{one_hot}_w, W)$$

²Tweaking this procedure could result in a language model

Skip-Gram

Embedding Computation

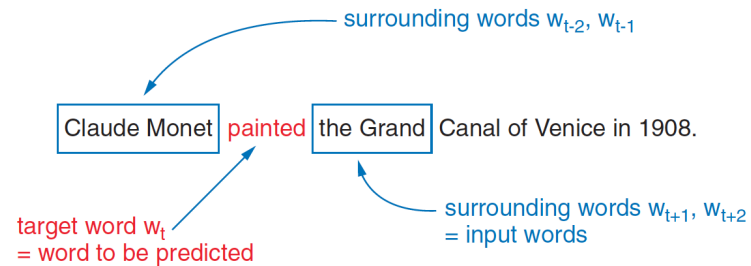


CBOW

Definition Continuous bag-of-words

Input context words

Output target (centre) word



(Lane et al., 2019, p. 196)

CBOW

Learning the Representations (1/3)

- Window size: 2 words \rightarrow 5-grams
- Input: multi-hot vector (sum of all one-hot vectors)
- Output: one-hot vector

$$S = w_1 w_2 w_3 w_4 w_5 w_6 w_7 w_8 w_9 w_{10}$$

$$[\dots] \underline{w_{t-2} w_{t-1}} w_t \underline{w_{t+1} w_{t+2}} [\dots]$$

CBOW

Learning the Representations (2/3)

Example: "Claude Monet painted the Grand Canal of Venice in 1908."

| input | | | | expected output |
|-----------|-----------|-----------|-----------|-----------------|
| w_{t-2} | w_{t-1} | w_{t+1} | w_{t+2} | w_t |
| | | Monet | painted | Claude |
| | Claude | painted | the | Monet |
| Claude | Monet | the | Grand | painted |
| Monet | painted | Grand | Canal | the |
| painted | the | Canal | of | Grand |
| the | Grand | of | Venice | Canal |
| Grand | Canal | Venice | in | of |
| Canal | of | in | 1908 | Venice |
| of | Venice | 1908 | | in |
| Venice | in | | | 1908 |


(Lane et al., 2019, p. 194)

CBOW

Learning the Representations (3/3)

Training

- The input is a multi-hot vector:
 $w_{t-2} + w_{t-1} + w_{t+2} + w_{t+2}$
- The output is a one-hot vector w_t



Final Comments

Skip-gram

- Works well with small corpora
- Some high-frequency [2, 3]-grams are added as single terms (e.g., New_York, Chicago_Bears)
- High-frequency tokens are subsampled (\sim to IDF over stopwords)
- Negative sampling. Not all weights are updated give a pair, just a few negative samples (much cheaper, roughly the same result)

CBOW

- Higher accuracy for frequent words
- Much faster to train

Next time

- Hands on word embeddings

References

- Lane, H., C. Howard, and H. Hapkem
2019. *Natural Language Processing in Action*. Shelter Island, NY: Manning Publication Co.
- Mikolov, T., K. Chen, G. Corrado, and J. Dean
2013. Efficient estimation of word representations in vector space. In *Arxiv*.